

DistriNVS: Distributional Geometry Transport and Visibility-Constrained Completion for Cross-Endoscope Novel-View Synthesis

Anonymous Authors

Abstract—Cross-endoscope novel-view synthesis must both transport tissue observed by the source camera and infer tissue revealed only from the target viewpoint. These two regimes are confounded when sparse or erroneous depth is projected as deterministic points: sampling gaps, depth failures, surface collisions, and genuine disocclusions all become visually identical holes. We present DistriNVS, a reliability-aware framework that models every source depth as a distribution and propagates it through calibrated geometry. Its distributional surface splatting renderer produces an anisotropic target-plane footprint together with support, depth variance, and collision entropy. These statistics drive a transport-or-synthesize router: a residual transport expert corrects co-visible appearance, whereas an uncertainty-conditioned Fourier network synthesizes unsupported or geometrically unreliable content. A final hard composition prevents the network from modifying trustworthy projected pixels. Joint target RGB-depth prediction, backward reprojection, and an explicit risk head provide geometric and failure-awareness constraints. The prescribed protocol trains on a fixed set of 13 iMED scenes and evaluates without adaptation on held-out EndoVis datasets 8 and 9 using a strict raw-overlap protocol.

I. INTRODUCTION

Robot-assisted minimally invasive surgery depends on endoscopic vision for scene perception and instrument guidance [1], [2]. Cross-endoscope novel-view synthesis can broaden this visual context by predicting a target-endoscope view from synchronized appearance and geometry acquired by a spatially displaced source endoscope [3]. Unlike left-to-right synthesis within one stereo endoscope, the target may observe anatomy absent from both source cameras. The task therefore combines two epistemically distinct regimes: co-visible anatomy has measured appearance that should be transported through calibrated geometry, whereas target-only anatomy lacks source evidence and must be synthesized. A reliable system should preserve supported measurements, identify inferred content, and thereby support wider-field robotic perception, retrospective visualization, and surgical simulation [4]. Such uncertainty-aware cross-view completion may also support visual-state construction for surgical robot world models [5].

Achieving this separation is difficult because endoscopic depth and calibration are unreliable around specularities, weak texture, depth discontinuities, instruments, and deforming tissue [6]. Conventional forward warping represents each valid source pixel by one 3-D point; discretized projection leaves sampling gaps, while multiple points may compete for the same target pixel. A binary validity mask then conflates




Fig. 1. Cross-endoscope synthesis with explicit provenance. DSS transports source-supported anatomy and exposes visibility uncertainty; DistriNVS synthesizes target-only content while trusted composition preserves reliable projected RGB-D.

missing depth, rasterization holes, depth-ordering conflicts, and genuine cross-endoscope disocclusions. Treating these cases uniformly either degrades observed structure or permits hallucination where geometric evidence remains trustworthy.

Existing methods address complementary parts of this problem but do not expose the required evidence boundary. Differentiable point rendering and softmax splatting improve feature transport and collision handling [7], [8], while reference-guided inpainting combines depth-based alignment with learned completion [9]. Large-mask inpainting networks provide strong global image priors [10], [11], and endoscopic video completion adds depth, temporal, or stereo cues [12]–[14]. Endoscopic neural rendering has likewise advanced sparse-view reconstruction and uncertainty modeling [15], [16]. However, deterministic projection commonly assigns equal reliability to all valid depths, while image-space uncertainty does not explicitly describe how source geometry supports a spatially offset target view. Neither approach propagates per-pixel depth uncertainty through the cross-endoscope transform and uses the resulting visibility statistics to constrain where appearance may be synthesized.

Our central premise is to *transport reliable geometric evidence and synthesize only unsupported or geometrically unreliable regions*. We introduce DistriNVS, a reliability-aware framework for direct cross-endoscope novel-view synthesis. Its distributional surface splatting (DSS) represents each source depth as a distribution whose bounded mean and scale, surface normal, and confidence are predicted by a depth reliability encoder. DSS projects deterministic sigma samples through the calibrated cross-endoscope transform and aggregates source features using Jacobian- and normal-aware footprints with soft nearest-surface visibility, yielding target support, depth variance, and collision entropy. These

statistics drive a continuous gate between residual transport for co-visible regions and uncertainty-conditioned Fourier completion for unsupported or ambiguous regions. A trusted mask hard-copies trusted DSS-supported RGB-D into the output, while target-depth and risk prediction plus backward RGB-D reprojection constrain completion during training.

The paper makes three contributions:

- We formulate cross-endoscope novel-view synthesis as a reliability-aware separation between transport of anatomy observed by the source endoscope and constrained inference of content revealed only by the target endoscope.
- We introduce DSS, a differentiable renderer that propagates depth distributions through calibrated cross-endoscope geometry and exposes target support, depth variance, and collision entropy beyond binary validity.
- We design a two-expert decoder that routes co-visible regions to residual transport and unsupported or geometrically unreliable regions to uncertainty-conditioned Fourier completion, while trusted-pixel composition preserves reliable projected RGB-D.

II. RELATED WORK

We briefly review image completion, geometry-guided view transport, and endoscopic neural rendering, with particular attention to how existing approaches represent missing content, exploit calibrated geometry, and quantify uncertainty.

A. Image and Endoscopic Video Inpainting

Generic image inpainting combines mask-aware feature extraction with increasingly broad appearance priors. Partial and gated convolutions prevent invalid activations from contaminating local features [17], [18], whereas LaMa and UFFC use Fourier operations for global context and MAT restricts attention to valid tokens [10], [11], [19]. Recent diffusion models such as BrushNet and TurboFill incorporate masked-image features into pretrained generators or adapt few-step sampling for efficient completion [20], [21]. Despite strong appearance priors, these methods define holes in a fixed image coordinate system and do not describe geometric provenance. Endoscopic completion adds domain cues: temporal learning removes specularities, DAEVI combines spatio-temporal RGB and depth, and SSIFNet fuses flow and stereo-temporal information for instrument occlusions [12]–[14]. Such methods exploit neighboring frames or views but still restore defects defined in the original camera coordinates. Cross-endoscope synthesis instead creates holes through viewpoint transfer: target-only anatomy is unobserved, while rasterization gaps or depth conflicts may retain source support. DistriNVS therefore uses calibrated projection statistics to distinguish true disocclusions from rasterization gaps and depth conflicts.

B. Geometry-Guided Transport and Completion

Geometry-guided view synthesis preserves correspondence by transporting source evidence before generating

missing content. SynSin predicts and projects a feature point cloud, softmax splatting resolves many-to-one mappings with learned importance, and GeoFill refines depth and pose before reprojection and inpainting [7]–[9]. Recent feed-forward Gaussian renderers use stronger explicit representations: MVSplat localizes Gaussians with plane-sweep cost volumes, whereas DepthSplat combines multi-view correspondence with monocular depth priors [22], [23]. They reduce per-scene optimization but infer a radiance representation from multiple calibrated observations, without an explicit per-pixel boundary between measured and synthesized appearance. Generative methods target wider gaps. LeftRefill frames reference-conditioned synthesis as paired-canvas completion, MultiDiff conditions video diffusion on depth-warped reference views, and point-guided latent diffusion uses stereo-derived clouds for laparoscopic view synthesis [24]–[26]. These models can generate plausible structures in large disocclusions, yet plausibility alone does not reveal which pixels remain geometrically supported. DistriNVS instead transports a calibrated source observation with distributional splatting, derives support, variance, and collision statistics during rasterization, and uses them to bound synthesis while preserving trusted pixels.

C. Endoscopic Novel-View Synthesis and Uncertainty

Endoscopic neural rendering must handle non-rigid tissue, view-dependent illumination, tool occlusions, and sparse camera coverage. EndoNeRF represents deformation with a neural radiance field, while EndoSurf combines dynamic modeling with a neural surface [27], [28]. EndoSparse introduces geometric and appearance priors into Gaussian splatting, and Endo-GSMT uses depth, dense-displacement, and motion priors for dynamic reconstruction [15], [29]. GastroNVS further shows that large viewpoint changes in real gastroscopy amplify geometric artifacts and inter-view illumination mismatch [30]. These systems build persistent scene representations from multiple frames and principally address reconstruction fidelity within a captured scene. Uncertainty-aware approaches address sparse evidence from another direction: UC-NeRF conditions a radiance field on multi-view uncertainty, whereas ExtraGS uses uncertainty to sample virtual cameras and weight diffusion-refined pseudo-observations for extrapolation [4], [16]. Thus, neither uncertainty nor generative priors are new to endoscopic rendering. The distinction is operational: DSS propagates a distribution for each input depth through the calibrated cross-endoscope transform and derives target support, depth variance, and collision entropy from the same rasterization that transports appearance. DistriNVS uses these quantities to route completion and protect trusted measurements during unseen-sequence inference without target-scene optimization.

III. METHOD

A. Problem Formulation and Overview

After resizing, both endoscope views have spatial size $H \times W$; we omit the batch dimension throughout. DistriNVS receives source RGB $\mathbf{I}_s \in [0, 1]^{3 \times H \times W}$, sparse metric depth




Fig. 2. DistriNVS framework. A reliability encoder predicts source geometry and features. DSS transports deterministic depth hypotheses through the calibrated cross-endoscope transform and jointly produces warped evidence and target-plane visibility statistics. These statistics condition a bounded transport-or-synthesize gate, while an independent trusted mask hard-copies accepted DSS evidence. Solid arrows denote inference; the dashed lower lane is used only during training.

$D_s \in \mathbb{R}_+^{1 \times H \times W}$, and validity mask $M_s \in \{0, 1\}^{1 \times H \times W}$. Calibration comprises intrinsics $K_s, K_t \in \mathbb{R}^{3 \times 3}$ and a rigid transform $T_{s \rightarrow t} \in \text{SE}(3) \subset \mathbb{R}^{4 \times 4}$. The model predicts target RGB $\hat{I}_t \in [0, 1]^{3 \times H \times W}$, depth $\hat{D}_t \in \mathbb{R}_+^{1 \times H \times W}$, and risk $U_t \in [0, 1]^{1 \times H \times W}$. Target RGB-D and the target instrument mask are used only for training or evaluation.

Figure 2 summarizes the inference pipeline. A reliability encoder estimates source geometry and uncertainty. Distributional surface splatting (DSS) then transports source evidence and measures target-view support and ambiguity. These measurements route complementary transport and synthesis experts. Finally, hard composition restores trusted rendered evidence exactly. The backward photometric cycle and risk objective are used only during training.

B. Reliability-Aware Source Geometry

Naive depth completion erases the distinction between measured and inferred geometry. The reliability encoder instead processes source RGB, median-normalized depth, and its validity mask. It produces a feature map $F_s \in \mathbb{R}^{d_f \times H \times W}$ and scalar fields $\mu, \sigma, c \in \mathbb{R}^{1 \times H \times W}$. They represent depth mean and scale in millimetres, and normalized confidence. It also predicts a unit orientation field $n \in \mathbb{R}^{3 \times H \times W}$. The normalization scale s_D is a positive scalar in millimetres, chosen as the median valid source depth, and missing inputs are set to zero. During training, M_s is the post-dropout input mask, whereas M_s^{gt} has the same dimensions and retains the pre-dropout supervision.

Measured depths receive a bounded residual correction; missing depths use a positive completion branch. Both obey the physical range, but disjoint confidence intervals keep inferred geometry less authoritative than measurements. The optical-axis-stabilized orientation proxy is learned only through downstream objectives. It serves as a footprint cue




Fig. 3. How continuous DSS support becomes a valid warp, illustrated on a zero-shot EndoVis dataset 8 result. (a) The completed-geometry DSS warp; unsupported pixels are black. (b) Unnormalized support, with colors capped at q_{vis} only for display; the cyan contour marks the validity threshold. (c) The exact render-validity mask is computed from uncapped support.

rather than a metrically supervised normal. Fixed ranges and clipping operations appear in Sec. IV-B.

C. Distributional Surface Splatting

DSS concatenates source RGB and features as $A_s = [I_s, F_s]$ and transports it to a target tensor A_t^w . Both tensors contain $3 + d_f$ channels over the $H \times W$ grid. DSS simultaneously predicts expected depth $\bar{D}_t \in \mathbb{R}_+^{1 \times H \times W}$. Support S , coverage C , depth variance V , collision entropy $\mathcal{H}$, confidence κ , and render validity M^{dss} are $1 \times H \times W$ target-grid maps. Support and variance are non-negative; coverage, collision entropy, and confidence lie in $[0, 1]$; validity is binary. All arise from one rasterization of a compact depth distribution at each source pixel. Expected depth is in millimetres and variance in square millimetres; the remaining evidence maps are dimensionless.

Depth hypotheses. For each source pixel q , depth hypothesis z_{qk} corresponds to a fixed standardized offset $\xi_k \in \Xi$ and quadrature weight a_k . The set and weights are specified in Sec. IV-B. The calibrated pinhole operator projects each hypothesis to a target coordinate $x_{qk} \in \mathbb{R}^2$ and target

depth $z_{qk}^t \in \mathbb{R}_+$. It comprises source back-projection, rigid transformation, and target projection [7], [8].

Uncertainty-aware visibility. Depth uncertainty is converted into a hypothesis-dependent, axis-aligned target footprint. The projection derivative J_{qk}^α measures the sensitivity of image coordinate $\alpha \in \{u, v\}$ to source depth. It has units of pixels per millimetre. The pixel-space footprint is

$$r_{qk}^\alpha = b + J_{qk}^\alpha \sigma_q + b_n(1 - |n_q^z|). \quad (1)$$

The radii r_{qk}^α , b , and b_n are measured in pixels, while n_q^z is the optical-axis component of the orientation proxy. For target pixel p , ρ_{qkp} denotes its elliptical distance from the projected hypothesis, normalized by both radii. The corresponding contribution is

$$w_{qk \rightarrow p} = c_q a_k \exp \left[-\frac{\rho_{qkp}^2}{2} - \frac{[z_{qk}^t - z_{\min}(p)]_+}{\tau_z} \right]. \quad (2)$$

Source confidence and quadrature mass scale this contribution. Spatial proximity and soft nearest-surface visibility form the exponential penalties. Here, $z_{\min}(p)$ is the detached nearest arrival, $[\cdot]_+$ is the positive part, and τ_z is the visibility temperature in millimetres. Nearest-depth selection and integer assignment are discrete; all remaining operations are differentiable.

Reliability statistics. Let Ω_p collect the retained source-hypothesis contributions to target pixel p . Their total weight defines raw support S_p . At render-valid pixels, DSS normalizes these weights over Ω_p . Their weighted moments yield the transported tensor, expected depth, and conditional depth variance. Coverage is a bounded saturation of raw support. The normalization constant and exact coverage transform are reported in Sec. IV-B.

DSS couples normalized weight entropy to depth dispersion. Let $H_N(p)$ denote the normalized Shannon entropy of the retained weights at pixel p . Collision ambiguity is

$$\mathcal{H}_p = H_N(p) [1 - \exp(-V_p/\tau_z^2)]. \quad (3)$$

This coupling discounts entropy caused only by benign sub-pixel overlap. We derive target confidence κ_p by penalizing coverage with depth variance and collision ambiguity. A pixel is render-valid when its raw support reaches $S_{\min}$, as shown in Fig. 3. Invalid pixels retain support and coverage but zero the transported outputs and ambiguity maps. The continuous evidence therefore distinguishes weak support from conflicting depth order rather than reducing reliability to a binary mask.

D. Reliability-Guided Completion and Preservation

Physics-guided routing. Both experts receive transported evidence and a condition tensor $\mathbf{h} \in \mathbb{R}^{6 \times H \times W}$. Its channels are coverage, confidence, normalized variance, collision entropy, normalized depth, and DSS validity. An analytic prior $g^0 \in [0, 1]^{1 \times H \times W}$ favours synthesis as geometric reliability decreases and equals one outside DSS support. The pre-trust synthesis probability is

$$g_p^* = \text{sigmoid}[\ell_p^0 + \eta_g \tanh R_\theta(\mathbf{h})_p]. \quad (4)$$

Here, ℓ^0 is the detached logit of the analytic prior, and the residual router R_θ outputs a scalar target-grid map. The factor η_g bounds its correction. Routing therefore retains a geometric initialization while adapting locally. The prior coefficients and normalization transforms are reported in Sec. IV-B. Both g^* and the effective gate g have shape $1 \times H \times W$.

Complementary experts. The shared expert input has $(9 + d_f) \times H \times W$ channels. The transport expert predicts a bounded RGB correction to the DSS warp, yielding $\mathbf{I}_t^{\text{trans}} \in [0, 1]^{3 \times H \times W}$. It has no depth head, so its depth remains the DSS expectation. The synthesis expert is an uncertainty-conditioned Fourier U-Net [31] that combines local convolution with real-FFT processing. Its outputs $(\mathbf{I}_t^{\text{syn}}, \mathbf{D}_t^{\text{syn}}, \mathbf{U}_t^{\text{syn}})$ have channel dimensions 3, 1, and 1 over the target grid. This local-global design follows Fourier inpainting networks [10], [11], but all weights are trained within DistriNVS. Risk remains an error-scale proxy, not a calibrated failure probability.

Trusted composition. Soft routing alone cannot guarantee preservation. The raw mask $M^{\text{trust},0} \in \{0, 1\}^{1 \times H \times W}$ requires sufficient support, low variance, low collision ambiguity, and DSS validity. Spatial cleanup removes isolated untrusted islands and yields M^{trust} with the same shape. The raw mask forces the effective gate g to zero. Elsewhere, g linearly blends the transport and synthesis RGB-D predictions. Let $\mathbf{Z}_t^{\text{mix}} \in \mathbb{R}^{5 \times H \times W}$ contain the blended RGB, depth, and synthesis risk. Let $\mathbf{Z}_t^w$ contain DSS RGB-D and a zero risk channel with the same dimensions. Hard preservation produces

$$\hat{\mathbf{Z}}_t = M^{\text{trust}} \odot \mathbf{Z}_t^w + (1 - M^{\text{trust}}) \odot \mathbf{Z}_t^{\text{mix}}. \quad (5)$$

Here, $\hat{\mathbf{Z}}_t \in \mathbb{R}^{5 \times H \times W}$ stacks the final RGB, depth, and risk, and $\odot$ broadcasts a mask across channels. Every trusted output equals the accepted DSS warp and has zero predicted risk. This zero-drift guarantee is relative to the warp, not to ground truth.

Inference-only repair. Reported runs additionally apply Telea fast-marching inpainting [32] to small, connected untrusted components. The recorded repair mask affects RGB only; larger components remain learned predictions, trusted pixels are never changed, and the repair is absent during training. Exact component and radius settings are reported in Sec. IV-B.

E. Learning Objectives

The training target $\mathbf{I}_t \in [0, 1]^{3 \times H \times W}$ is accompanied by a tissue-valid mask $M_{\text{tis}} \in \{0, 1\}^{1 \times H \times W}$, which excludes annotated instruments when available. The untrusted, trusted, and seam regions M_u , M_v , and M_b are single-channel masks derived from M^{trust} and share its dimensions. Their Charbonnier terms compare predicted and target RGB; $\mathcal{L}_\nabla$ additionally matches image gradients near untrusted regions and seams. These training regions differ from the raw-warp masks used for evaluation.

A masked Laplace negative log-likelihood $\mathcal{L}_{d-\text{nl}}$ supervises the source depth mean and scale on the pre-dropout

mask M_s^{gt} . Available target depth D_t shares the dimensions of $\tilde{D}_t$ and supplies a masked relative-depth term $\mathcal{L}_{t\text{-depth}}$. The depth-guided cycle splats the predicted RGB-D back through the inverse calibrated transform. Detached risk controls its depth scale, whereas complementary risk controls its confidence. The cycle term $\mathcal{L}_{\text{cyc}}$ compares backward-rendered RGB with the source image where both are valid. We do not apply a depth-cycle loss.

A binary router target marks every pixel that violates any trust test, and $\mathcal{L}_{\text{router}}$ applies binary cross-entropy with logits over the full image. For a mask M , $\langle \cdot \rangle_M$ denotes the mean over its selected entries. The scalar e_p is the channel-mean absolute RGB error at pixel p . The positive scale s_p is an affine mapping of predicted risk between fixed bounds. The uncertainty objective is

$$\mathcal{L}_u = \langle e_p / s_p + \log s_p \rangle_{M_u}. \quad (6)$$

Exact risk bounds are given in Sec. IV-B. The appearance group $\mathcal{L}_{\text{app}}$ combines the regional RGB terms and the gradient term. The depth group $\mathcal{L}_{\text{depth}}$ combines source likelihood and target relative-depth supervision. The target term is omitted when target depth is unavailable. The complete objective is

$$\mathcal{L} = \mathcal{L}_{\text{app}} + \lambda_d \mathcal{L}_{\text{depth}} + \lambda_c \mathcal{L}_{\text{cyc}} + \lambda_r \mathcal{L}_{\text{router}} + \lambda_u \mathcal{L}_u. \quad (7)$$

All loss weights are non-negative and are specified in the experimental setup. All losses and their weights are scalars.

At inference, the learned path ends at Eq. (5), followed in reported runs only by the disclosed deterministic repair. Training targets, tool masks, and the backward-cycle branch are not inference inputs. The method assumes known zero-skew pinhole intrinsics, a rigid inter-endoscope transform, and an axis-aligned approximation to projected uncertainty; these assumptions delimit its present scope.

IV. EXPERIMENTS

A. Datasets and Zero-Shot Protocol

Training uses the frozen list of 13 iMED Task 2 scene directories supplied in the project configuration [33]. Splits are made at scene level, not at frame level. The training loader uses a synchronized source RGB-depth and target RGB pair at 640×512 resolution. Fifteen percent of valid source depth samples are removed online, with additional 9×9 missing islands, so that the reliability encoder observes both isolated and contiguous depth failures. EndoVis/SCARED datasets 8 and 9 are reserved for final zero-shot evaluation; neither their images nor metrics are used for model selection.

We determine inclusion using a renderer-only pass from the *raw* valid source depth. A frame is evaluated only if $|M_{\text{raw}}^{\text{dss}}| / (H_I W_I) > 0.5$; the learned depth completion, target image, and target tool mask cannot influence this decision. This reproduces the project’s requested strict overlap criterion and prevents a learned module from selecting favourable frames. The current tables report frame-wise means and sample standard deviations, both rounded to two decimal places. Paired sequence-level confidence intervals and permutation tests will be added once all baseline predictions are frozen.

[TODO: Verify the final dataset provenance, license statement, total training pair count, held-out sequence count, and the immutable list of selected test frames before submission.]

B. Baselines, Metrics, and Implementation

Geometry baselines will include hard point projection, nearest-neighbour small-hole filling, deterministic soft splatting, and softmax splatting. Image completion baselines will include Telea, LaMa, UFFC, MAT, Stable Diffusion 1.5 Inpainting, and the existing LCM+endoscopy-LoRA system. Directly relevant endoscopic baselines (DAEVI and SSIFNet) will be evaluated when code and an equivalent cross-view protocol are available; otherwise their non-comparable reported values will not be inserted into the main table. Every completion method will use the same raw mask, spatial resolution, selected frames, and known-pixel hard composition.

The existing evaluator reports full-, raw-hole-, raw-visible-, and morphological seam-band PSNR/SSIM, full-image and hole-region LPIPS, known-region drift, risk expected calibration error (ECE), area under the risk-coverage curve (AURC), and a projection-defect taxonomy. The final study will additionally report DISTs, depth AbsRel/RMSE, disparity endpoint error, normal error, latency, throughput, peak memory, and parameter count. Results will be stratified by raw overlap, hole ratio, depth validity, instrument presence, and sequence.

Renderer and reliability. We use $\Xi = (-1, 0, 1)$ and $(a_k) = (0.2741, 0.4518, 0.2741)$; the display percentile in Fig. 3 is $q_{\text{vis}} = 99$ and affects visualization only. Before projection, every sampled depth hypothesis z_{qk} is clipped to $[1, 500]$ mm. Hypotheses affect the radius-one integer neighbourhood, with $\tau_z = 2$ mm and $S_{\text{min}} = 0.03$. Each footprint r_{qk}^α is clipped to $[0.35, 2.5]$ px, with $b = b_n = 0.75$ px. Spatial and final weights no greater than 10^{-5} and 10^{-8} are pruned, and $\epsilon_{\text{dss}} = 10^{-8}$ stabilizes normalization. Coverage uses $C = 1 - \exp(-S)$, and target confidence uses $\kappa = C \exp[-V / (4\tau_z^2) - \mathcal{H}]$. Weight entropy is divided by the logarithm of at least two retained contributions before it is coupled to depth variance.

The reliability encoder fuses full-resolution residual and quarter-resolution context branches to output $d_f = 16$ feature channels. Its fallback median is $s_D = 100$ mm, and valid depth is normalized as $\tilde{D}_s = \text{clip}(D_s / s_D, 0, 5)$ before missing values are set to zero. Measured and completed confidence occupy $[0.5, 1]$ and $[0, 0.35]$, respectively. The measured-depth correction ratio is bounded by 0.1, whereas the completion branch has a 0.05 positive offset. Depth and uncertainty are clipped to $[1, 500]$ mm and $[0.25, 20]$ mm.

Router and experts. The condition maps use normalized variance $\bar{V} = V / (V + \tau_z^2)$ and normalized depth $\tilde{D} = \text{clip}(\bar{D}_t / s_D, 0, 5) / 5$. The analytic prior is clipped to $[0, 1]$ and then to $[\epsilon_g, 1 - \epsilon_g]$ before its detached logit ℓ^0 is evaluated. For Eq. (4), the prior weights are $(\alpha_\kappa, \alpha_V, \alpha_H) = (0.65, 0.20, 0.15)$, $\epsilon_g = 10^{-4}$, and $\eta_g = 0.25$; the residual router is zero-initialized. The transport RGB residual is

bounded by 0.15. The four-resolution synthesis network uses 3×3 local branches and five output channels; its metric depth is $\lceil \text{softplus}(\cdot) + 0.05 \rceil s_D$. The shared base width is 32.

Masks, losses, and repair. The raw trust tests use $S_{\text{trust}} = 0.35$, $V_{\text{trust}} = 9 \text{ mm}^2$, and $\mathcal{H}_{\text{trust}} = 1.0$; cleanup uses a 7×7 opening kernel. Since collision entropy lies in $[0, 1]$, its configured hard threshold is neutral, although entropy still affects target confidence and soft routing. A separate 7×7 dilation-erosion kernel extracts the training seam. The Charbonnier constant is 10^{-3} , the source-depth likelihood uses a 10^{-3} scale floor, and target relative-depth error uses a 1-mm denominator floor. Equation (6) uses $(s_{\min}, s_{\max}) = (0.01, 0.50)$ in normalized-RGB units. For the backward cycle, the depth scale is $\text{clip}_{[0.25, 20]}[0.25 + 20 \text{sg}(\mathbf{U}_t)]$ mm and its confidence is $\text{clip}_{[0.05, 1]}(1 - \mathbf{U}_t)$. Reported inference applies radius-3 Telea repair to 8-connected untrusted components occupying at most 2% of the image.

Optimization. Training is scheduled for 10,000 AdamW steps with learning rate 2×10^{-4} , weight decay 10^{-4} , and three-GPU distributed training. The per-GPU batch size is 1 with four-step gradient accumulation, giving an effective batch size of 12. We use BF16 mixed precision, maximum gradient norm 1, and seed 6666. The implemented loss weights $(\lambda_h, \lambda_v, \lambda_b, \lambda_{\nabla}, \lambda_d, \lambda_c, \lambda_r, \lambda_u)$ are $(3, 1, 2, 0.2, 0.2, 0.2, 0.1, 0.1)$.

[TODO: Record the final checkpoint, model-selection rule, GPU model, training time, inference warm-up, precision, exact parameter/FLOP counts, and all baseline versions and checkpoints.]

C. Main Comparison

We evaluated every available prediction on the seven-scene iMED validation split. Complete-coverage methods produced predictions for all 1,392 frames, and every included prediction was re-evaluated against the same target image and raw-warp hole mask. The mask covered 25.84% of an image on average. Five Gaussian pipelines covered all seven scenes. The four published reconstruction baselines used dataset-provided ground-truth depth, whereas the additional StereoSurGS pipeline combined SurgicalGS with FoundationStereo-estimated depth [34], [35]. They are reference reconstruction baselines rather than matched feed-forward completion models.

PR-ENDO, EndoGS, and EndoGaussian met the original run-validity criterion on 1/7, 5/7, and 6/7 scenes (199, 994, and 1,193 frames), respectively. We report these method-dependent subsets as diagnostics, but exclude them from ranking and best-value selection. The EndoGaussian diagnostic retains all frames from its six accepted scenes, including 155 frames flagged as degenerate, rather than selectively removing failed frames. D4Recon yielded only degenerate renders on all 1,392 selected frames, while the remaining zero-success runs provided no evaluable render set; assigning either group an image-quality score would therefore be artificial. Table I reports the frame-wise mean and sample standard deviation under this protocol. The deterministic small-component repair affected $1.17 \pm 0.94\%$ of image

pixels per validation frame on average; the 2% criterion applies to each connected component, so their union may be larger.

Among methods with complete 7/7-scene coverage, DistriNVS achieved the best distortion- and structure-based result, improving full-image PSNR by 0.94 dB and SSIM by 0.05 over the strongest corresponding image-completion baselines. When the complete Gaussian reconstruction baselines are included, the SSIM margin narrows to 0.02 over Free-SurGS, whereas the PSNR margin remains 0.94 dB over the LCM baseline. The advantage was larger inside the missing region: hole PSNR increased by 3.57 dB over LaMa, while hole SSIM increased by 0.08 over the supervised LoRA baseline. However, LaMa obtained lower full-image LPIPS (0.47 ± 0.03 versus 0.55 ± 0.03) and lower hole LPIPS (0.48 ± 0.04 versus 0.52 ± 0.03). The results therefore show stronger pixel and structural fidelity for DistriNVS, but not a uniform advantage in learned perceptual distance. The partial-coverage diagnostics characterize different subsets and do not support pairwise superiority claims.

We separately evaluated DistriNVS on held-out EndoVis Datasets 8 and 9 using the strict raw overlap > 0.5 inclusion rule. Among 124 candidate frames, 78 met the rule. Table II reports the two datasets separately and jointly.

Across the 78 accepted frames, DistriNVS obtained 23.79 ± 2.49 dB PSNR, 0.68 ± 0.12 SSIM, 0.23 ± 0.06 LPIPS, 21.82 ± 3.81 dB hole PSNR, 0.61 ± 0.15 hole SSIM, and 0.28 ± 0.09 hole LPIPS. The deterministic repair affected $3.75 \pm 2.11\%$ of pixels per accepted frame. Known-region drift was zero under the model-defined trusted mask. Baseline inference on Datasets 8 and 9 is not yet complete; consequently, Table II does not support a cross-method zero-shot ranking. Paired confidence intervals and significance tests will be added after the final prediction set is frozen.

D. Ablations and Geometric Diagnostics

The core ablation in Table III follows the implementation flags. The hard-point variant disables depth completion, distributional samples, normal-dependent footprint, physics routing, and the backward cycle. The soft splat variant retains a deterministic Gaussian footprint. Subsequent rows add DSS, the physics router, hard composition, and the RGB-D-defined backward cycle. This ordering tests whether an improvement originates from ordinary anti-aliased splatting or from the proposed distributional visibility model.

[TODO: Add columns for seam error, visible-region PSNR, known drift, depth error, AURC, and runtime. Run focused DSS ablations for 1/3/5 sigma points, isotropic versus anisotropic footprints, normal removal, hard versus soft visibility, and individual $S/V/\mathcal{H}$ router inputs.]

[TODO: Report the paired with/without small-component Telea repair ablation; keep the repair disabled in all architecture-only ablations.]

To determine whether DSS solves the intended geometric problem before learned completion is considered, we will report raw coverage, visible-region RGB error, target-depth

TABLE I

COMPARISON ON THE SEVEN-SCENE IMED VALIDATION SPLIT. ROWS WITH $N = 1,392$ COVER ALL SEVEN SCENES; PARTIAL ROWS SUMMARIZE ONLY THEIR ACCEPTED SCENES AND ARE DIAGNOSTIC. EVERY INCLUDED PREDICTION USES THE SAME TARGET AND RAW-WARP HOLE MASK. VALUES ARE FRAME-WISE MEAN $\pm$ SAMPLE SD. BOLD DENOTES THE BEST MEAN AMONG COMPLETE-COVERAGE ROWS ONLY. FS DENOTES FOUNDATIONSTEREO-ESTIMATED DEPTH; THE OTHER GAUSSIAN RUNS USE DATASET GROUND-TRUTH DEPTH. THE DISTRINVS ROW INCLUDES THE DISCLOSED SMALL-COMPONENT TELEA REPAIR.

Method	N	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	H-PSNR $\uparrow$	H-SSIM $\uparrow$	H-LPIPS $\downarrow$
<i>Gaussian reconstruction—complete coverage</i>							
Deform3DGS [36]	1,392	15.55 $\pm$ 1.78	0.55 $\pm$ 0.04	0.63 $\pm$ 0.06	14.11 $\pm$ 1.37	0.47 $\pm$ 0.10	0.69 $\pm$ 0.06
Endo-4DGS [37]	1,392	16.44 $\pm$ 1.71	0.58 $\pm$ 0.04	0.61 $\pm$ 0.06	15.61 $\pm$ 1.47	0.57 $\pm$ 0.13	0.66 $\pm$ 0.06
Free-SurGS [38]	1,392	16.83 $\pm$ 2.11	0.67 $\pm$ 0.06	0.63 $\pm$ 0.04	17.05 $\pm$ 2.38	0.72 $\pm$ 0.04	0.61 $\pm$ 0.04
SurgicalGS [34]	1,392	17.41 $\pm$ 1.99	0.60 $\pm$ 0.06	0.60 $\pm$ 0.05	16.98 $\pm$ 2.55	0.61 $\pm$ 0.14	0.65 $\pm$ 0.04
StereoSurGS (FS)	1,392	16.37 $\pm$ 1.46	0.52 $\pm$ 0.08	0.59 $\pm$ 0.05	13.59 $\pm$ 1.27	0.39 $\pm$ 0.13	0.70 $\pm$ 0.06
<i>Gaussian reconstruction—partial coverage (diagnostic)</i>							
PR-ENDO [39]	199	15.51 $\pm$ 0.76	0.50 $\pm$ 0.01	0.62 $\pm$ 0.01	12.98 $\pm$ 0.46	0.20 $\pm$ 0.05	0.71 $\pm$ 0.02
EndoGS [40]	994	15.86 $\pm$ 2.19	0.61 $\pm$ 0.06	0.71 $\pm$ 0.11	15.50 $\pm$ 1.97	0.65 $\pm$ 0.08	0.70 $\pm$ 0.07
EndoGaussian [41]	1,193	15.21 $\pm$ 3.56	0.57 $\pm$ 0.26	0.81 $\pm$ 0.10	16.00 $\pm$ 3.15	0.62 $\pm$ 0.28	0.72 $\pm$ 0.08
<i>Image completion</i>							
SD1.5 Inpainting	1,392	16.84 $\pm$ 2.73	0.57 $\pm$ 0.05	0.52 $\pm$ 0.06	14.56 $\pm$ 3.14	0.50 $\pm$ 0.10	0.64 $\pm$ 0.07
LaMa	1,392	18.99 $\pm$ 1.58	0.64 $\pm$ 0.07	0.47 $\pm$ 0.03	20.41 $\pm$ 2.06	0.74 $\pm$ 0.05	0.48 $\pm$ 0.04
LCM + endoscopy LoRA	1,392	19.02 $\pm$ 1.64	0.63 $\pm$ 0.09	0.49 $\pm$ 0.03	19.54 $\pm$ 2.28	0.74 $\pm$ 0.09	0.49 $\pm$ 0.03
DistriNVS (ours)	1,392	19.96 $\pm$ 1.79	0.69 $\pm$ 0.08	0.55 $\pm$ 0.03	23.98 $\pm$ 3.28	0.82 $\pm$ 0.06	0.52 $\pm$ 0.03

TABLE II

DISTRINVS ZERO-SHOT PERFORMANCE ON HELD-OUT ENDOVIS DATASETS 8 AND 9. VALUES ARE FRAME-WISE MEAN $\pm$ SD OVER FRAMES SATISFYING raw overlap > 0.5 AND INCLUDE THE DISCLOSED SMALL-COMPONENT REPAIR.

Dataset	Frames	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	H-PSNR $\uparrow$	H-SSIM $\uparrow$	H-LPIPS $\downarrow$
EndoVis 8	48	24.85 $\pm$ 2.00	0.72 $\pm$ 0.12	0.21 $\pm$ 0.04	23.71 $\pm$ 2.74	0.70 $\pm$ 0.09	0.23 $\pm$ 0.06
EndoVis 9	30	22.10 $\pm$ 2.29	0.63 $\pm$ 0.09	0.27 $\pm$ 0.06	18.80 $\pm$ 3.32	0.45 $\pm$ 0.08	0.35 $\pm$ 0.08
Combined	78	23.79 $\pm$ 2.49	0.68 $\pm$ 0.12	0.23 $\pm$ 0.06	21.82 $\pm$ 3.81	0.61 $\pm$ 0.15	0.28 $\pm$ 0.09

TABLE III

COMPONENT ABLATION. K IS THE NUMBER OF DEPTH SIGMA POINTS.

Variant	K	DSS	Router	Hard	PSNR
Hard point	1	—	—	✓	—
Soft splat	1	—	—	✓	—
DSS only	3	✓	—	✓	—
DSS + router	3	✓	✓	✓	—
No hard composition	3	✓	✓	—	—
No backward cycle	3	✓	✓	✓	—
Full DistriNVS	3	✓	✓	✓	—

error, and occlusion-boundary ordering for hard, softmax, and distributional splatting. The diagnostic label map partitions target pixels into reliable/unclassified support, local sampling gaps, depth-completion-recoverable gaps, unsupported/disoccluded regions, and surface-collision risk. These labels are diagnostics rather than ground truth because a single source view cannot perfectly separate genuine disocclusion from catastrophic depth failure.

[TODO: Insert the DSS-only table and plots showing monotonic relationships between support/variance/collision entropy and actual reprojection error.]

E. Qualitative, Failure, and Efficiency Analysis

Efficiency is central to the ICRA framing. We will measure end-to-end latency including depth reliability estimation, DSS, both experts, and hard composition; DSS-only timing will also be reported to isolate the renderer cost. Failure analysis will cover specular depth loss, calibration perturbation, thin instruments, large disocclusion, weak texture, and dynamic tissue deformation.

[TODO: Insert the speed/memory/parameter table, calibration-perturbation curves, and blinded expert assessment if available. Do not claim real-time performance until measured on the stated hardware and resolution.]

V. DISCUSSION

Under a common validation protocol, DistriNVS outperformed the evaluated Gaussian reconstruction, generic inpainting, and supervised LoRA baselines in both full-image and hole-region distortion and structural fidelity. The Gaussian rows provide a reference rather than a matched training comparison, because they are optimized per scene with dataset ground-truth depth. The larger improvement in hole PSNR (3.57 dB) than in full-image PSNR (0.94 dB) is consistent with the intended separation between evidence transport and unsupported-region synthesis. Together with

[TODO: Replace this placeholder with synchronized columns: source RGB, raw hard warp, raw validity, LaMa, LCM, DistriNVS, target RGB, support, variance, collision entropy, and risk. Include at least one success and one failure for each overlap bin.]

Fig. 4. Planned qualitative comparison and failure visualization. Crops and color scales must be identical across methods; generated and source-backed regions will be explicitly marked.

zero measured drift under the trusted mask, these observations support the operational value of constraining where synthesis is permitted. They do not yet isolate the individual contributions of distributional splatting, routing, and hard composition; the planned component ablations remain necessary for that causal claim. Moreover, DistriNVS’s higher LPIPS than LaMa shows that its gains in pixel and structural agreement do not translate uniformly to the AlexNet perceptual feature space; qualitative analysis is needed to determine whether this difference reflects texture realism, colour statistics, or metric bias.

The method does not remove the fundamental ambiguity of unseen tissue. A low risk score indicates model confidence under the training distribution; it does not certify anatomical truth. Moreover, the current formulation assumes known intrinsics and a rigid inter-camera transform for each pair, processes frames independently, and uses an axis-aligned rather than full-covariance footprint. The zero-shot evaluation also revealed a gap between the two held-out datasets: full-image PSNR decreased from 24.85 ± 2.00 dB on EndoVis 8 to 22.10 ± 2.29 dB on EndoVis 9, while hole SSIM decreased from 0.70 ± 0.09 to 0.45 ± 0.08 . This difference limits claims of uniform cross-dataset generalization, and incomplete baseline runs currently prevent a fair cross-method ranking on Datasets 8 and 9. Large calibration errors, non-rigid motion between nominally synchronized frames, or out-of-distribution anatomy may therefore invalidate the transport statistics. Generated output should remain visibly marked and should not be used as a substitute for a live physical endoscopic view.

[TODO: Rewrite the first discussion paragraph from the final ablation and risk curves. Add observed failure modes, not generic limitations, and compare them directly with LaMa/UFFC, DAEVI/SSIFNet, UC-NeRF, and ExtraGS.]

VI. CONCLUSION

This paper introduces a reliability-aware formulation of cross-endoscope novel-view synthesis and implements it through distributional surface splatting, visibility-constrained transport/synthesis routing, trusted-pixel hard composition, and joint appearance–geometry risk modeling. The design turns projection support, depth variance, and collision entropy into operational constraints on what the network

may synthesize. [TODO: After the frozen evaluation, replace the following validation evidence with the frozen test-set comparison if available.] On a common 1,392-frame validation protocol, DistriNVS exceeded the strongest complete-coverage per-metric baselines by 0.94 dB in full-image PSNR and 3.57 dB in hole PSNR, while attaining a hole SSIM of 0.82 ± 0.06 . These results support visibility-constrained completion under the evaluated setting, although LaMa retained the lowest LPIPS. The difference between EndoVis Datasets 8 and 9 and unfinished zero-shot baseline evaluation define the current boundary of the evidence.

REFERENCES

- [1] R. Wei, K. Chen, Y. L. Ng, Y. Ma, J. D.-L. Ho, H. S. Tong, X. Wang, J. Dai, K.-W. Kwok, and Q. Dou, “Real-time monocular 2-d and 3-d perception of endoluminal scenes for controlling flexible robotic endoscopic instruments,” *IEEE Transactions on Robotics*, vol. 42, pp. 1535–1555, 2026.
- [2] I. O. Sunmola, Z. Zhao, S. Schmidgall, Y. Wang, P. M. Scheikl, V. Pham, and A. Krieger, “Surgical gaussian surfels: Highly accurate real-time surgical scene rendering using gaussian surfels,” in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2026, pp. 4515–4524.
- [3] X. Wang, B. Yang, M. Wei, Y. Chen, J. Zhang, and Y. Nie, “Endo-E2E-GS: End-to-end 3d reconstruction of endoscopic scenes using gaussian splatting,” *Displays*, vol. 92, p. 103353, 2026.
- [4] C.-T. Hsieh, J. Shan, H. Fang, J. Hu, T. Ni, L. Han, Y. Ban, S. S. Cheng, and H. Wang, “ExtraGS: Enhancing endoscopic view extrapolation via diffusion-guided 3d gaussian splatting,” *arXiv preprint arXiv:2607.12785*, 2026.
- [5] Y. He, P. Guo, M. Xu, Z. Li, A. Myronenko, D. Imans, B. Liu, D. Yang, M. Gu, Y. Ji, Y. Jin, R. Zhao, B. Shen, and D. Xu, “SurgWorld: Learning surgical robot policies from videos via world modeling,” *arXiv preprint arXiv:2512.23162*, 2025. [Online]. Available: <https://arxiv.org/abs/2512.23162v1>
- [6] M. Allan, J. Mcleod, C. Wang, J. C. Rosenthal, Z. Hu, N. Gard, P. Eisert, K. X. Fu, T. Zeffiro, W. Xia, Z. Zhu, H. Luo, F. Jia, X. Zhang, X. Li, L. Sharan, T. Kurmann, S. Schmid, R. Sznitman, D. Psychogyios, M. Azizian, D. Stoyanov, L. Maier-Hein, and S. Speidel, “Stereo correspondence and reconstruction of endoscopic data challenge,” *arXiv preprint arXiv:2101.01133*, 2021.
- [7] O. Wiles, G. Gkioxari, R. Szeliski, and J. Johnson, “SynSin: End-to-end view synthesis from a single image,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 7467–7477.
- [8] S. Niklaus and F. Liu, “Softmax splatting for video frame interpolation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 5437–5446.
- [9] Y. Zhao, C. Barnes, Y. Zhou, E. Shechtman, S. Amirghodsi, and C. C. Fowlkes, “GeoFill: Reference-based image inpainting with better geometric understanding,” in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2023, pp. 1776–1786.

- [10] R. Suvorov, E. Logacheva, A. Mashikhin, A. Remizova, A. Ashukha, A. Silvestrov, N. Kong, H. Goka, K. Park, and V. Lempitsky, "Resolution-robust large mask inpainting with fourier convolutions," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2022, pp. 2149–2159.
- [11] T. Chu, J. Chen, J. Sun, S. Lian, Z. Wang, Z. Zuo, L. Zhao, W. Xing, and D. Lu, "Rethinking fast fourier convolution in image inpainting," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023, pp. 23 195–23 205.
- [12] R. Daher, F. Vasconcelos, and D. Stoyanov, "A temporal learning approach to inpainting endoscopic specularities and its effect on image correspondence," *Medical Image Analysis*, vol. 90, p. 102994, 2023.
- [13] F. X. Zhang, S. Chen, X. Xie, and H. P. H. Shum, "Depth-aware endoscopic video inpainting," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*, ser. Lecture Notes in Computer Science, vol. 15006. Springer Nature Switzerland, 2024, pp. 143–153.
- [14] X. Zou, Z. Zhang, D. Yu, W. Sun, W. Liu, D. Hang, W. Bao, and G. Zheng, "SSIFNet: Spatial-temporal stereo information fusion network for self-supervised surgical video inpainting," *Computerized Medical Imaging and Graphics*, vol. 125, p. 102622, 2025.
- [15] C. Li, B. Y. Feng, Y. Liu, H. Liu, C. Wang, W. Yu, and Y. Yuan, "EndoSparse: Real-time sparse view synthesis of endoscopic scenes using gaussian splatting," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*, ser. Lecture Notes in Computer Science, vol. 15006. Springer Nature Switzerland, 2024, pp. 252–262.
- [16] J. Guo, J. Wang, R. Wei, D. Kang, Q. Dou, and Y.-H. Liu, "UC-NeRF: Uncertainty-aware conditional neural radiance fields from endoscopic sparse views," *IEEE Transactions on Medical Imaging*, vol. 44, no. 3, pp. 1284–1296, 2025.
- [17] G. Liu, F. A. Reda, K. J. Shih, T.-C. Wang, A. Tao, and B. Catanzaro, "Image inpainting for irregular holes using partial convolutions," in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018, pp. 85–100.
- [18] J. Yu, Z. Lin, J. Yang, X. Shen, X. Lu, and T. S. Huang, "Free-form image inpainting with gated convolution," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019, pp. 4471–4480.
- [19] W. Li, Z. Lin, K. Zhou, L. Qi, Y. Wang, and J. Jia, "MAT: Mask-aware transformer for large hole image inpainting," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 10 758–10 768.
- [20] X. Ju, X. Liu, X. Wang, Y. Bian, Y. Shan, and Q. Xu, "BrushNet: A plug-and-play image inpainting model with decomposed dual-branch diffusion," in *Computer Vision – ECCV 2024*. Springer Nature Switzerland, 2024, pp. 150–168.
- [21] L. Xie, D. Pakhomov, Z. Wang, Z. Wu, Z. Chen, Y. Zhou, H. Zheng, Z. Zhang, Z. Lin, J. Zhou, and C. Dong, "TurboFill: Adapting few-step text-to-image model for fast image inpainting," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025, pp. 7613–7622.
- [22] Y. Chen, H. Xu, C. Zheng, B. Zhuang, M. Pollefeys, A. Geiger, T.-J. Cham, and J. Cai, "MVSplat: Efficient 3d gaussian splatting from sparse multi-view images," in *Computer Vision – ECCV 2024*, ser. Lecture Notes in Computer Science, vol. 15079. Springer Nature Switzerland, 2024, pp. 370–386.
- [23] H. Xu, S. Peng, F. Wang, H. Blum, D. Barath, A. Geiger, and M. Pollefeys, "DepthSplat: Connecting gaussian splatting and depth," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025, pp. 16 453–16 463.
- [24] C. Cao, Y. Cai, Q. Dong, Y. Wang, and Y. Fu, "LeftRefill: Filling right canvas based on left reference through generalized text-to-image diffusion model," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 7705–7715.
- [25] N. Müller, K. Schwarz, B. Rössle, L. Porzi, S. R. Bulò, M. Nießner, and P. Kotschieder, "MultiDiff: Consistent novel view synthesis from a single image," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 10 258–10 268.
- [26] W. Tang, T. Chen, Y. Espinel, S. Farid, E. Buc, A. Bartoli, and S. Ali, "Point-guided latent diffusion model for novel view synthesis in laparoscopic liver surgery," *Healthcare Technology Letters*, vol. 12, no. 1, p. e70032, 2025.
- [27] Y. Wang, Y. Long, S. H. Fan, and Q. Dou, "Neural rendering for stereo 3d reconstruction of deformable tissues in robotic surgery," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2022*. Springer, 2022, pp. 431–441.
- [28] R. Zha, X. Cheng, H. Li, M. Harandi, and Z. Ge, "EndoSurf: Neural surface reconstruction of deformable tissues with stereo endoscope videos," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2023*. Springer, 2023, pp. 13–23.
- [29] H. Gou, C. Wang, J. Yang, Y. Liu, F. Jia, D. Xiao, F. Qin, and H. Luo, "Endo-GSMT: Endoscopic monocular scene reconstruction with dynamic gaussian splatting and motion tracking," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2025*, ser. Lecture Notes in Computer Science, vol. 15968. Springer Nature Switzerland, 2025, pp. 213–223.
- [30] M. Minai, Y. Monno, M. Okutomi, and S. Suzuki, "Gastroendoscopy view synthesis: A new real dataset and evaluation," *arXiv preprint arXiv:2606.25427*, 2026.
- [31] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*. Springer International Publishing, 2015, pp. 234–241.
- [32] A. Telea, "An image inpainting technique based on the fast marching method," *Journal of Graphics Tools*, vol. 9, no. 1, pp. 23–34, 2004.
- [33] The iMED Challenge Organizers, "iMED Challenge 2026," MICCAI 2026 EndoVis Challenge, 2026, accessed: 2026-09-10. [Online]. Available: <https://imed-challenge.github.io/>
- [34] J. Chen, X. Zhang, M. I. Hoque, F. Vasconcelos, D. Stoyanov, D. S. Elson, and B. Huang, "SurgicalGS: Dynamic 3D Gaussian Splatting for Accurate Robotic-Assisted Surgical Scene Reconstruction," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2025*. Springer Nature Switzerland, 2025, pp. 572–582.
- [35] B. Wen, M. Trepte, J. Aribido, J. Kautz, O. Gallo, and S. Birchfield, "FoundationStereo: Zero-Shot Stereo Matching," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025, pp. 5249–5260.
- [36] S. Yang, Q. Li, D. Shen, B. Gong, Q. Dou, and Y. Jin, "Deform3DGS: Flexible Deformation for Fast Surgical Scene Reconstruction with Gaussian Splatting," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*. Springer Nature Switzerland, 2024, pp. 132–142.
- [37] Y. Huang, B. Cui, L. Bai, Z. Guo, M. Xu, M. Islam, and H. Ren, "Endo-4DGS: Endoscopic Monocular Scene Reconstruction with 4D Gaussian Splatting," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*. Springer Nature Switzerland, 2024, pp. 197–207.
- [38] J. Guo, J. Wang, D. Kang, W. Dong, W. Wang, and Y.-H. Liu, "Free-SurGS: SfM-Free 3D Gaussian Splatting for Surgical Scene Reconstruction," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*. Springer Nature Switzerland, 2024, pp. 350–360.
- [39] J. Kaleta, W. Smolak-Dyżewska, D. Malarz, D. Dall'Alba, P. Korzeniowski, and P. Spurek, "PR-ENDO: Physically Based Relightable Gaussian Splatting for Endoscopy," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2025*, vol. LNCS 15969. Springer Nature Switzerland, 2025, pp. 391–401.
- [40] L. Zhu, Z. Wang, J. Cui, Z. Jin, G. Lin, and L. Yu, "EndoGS: Deformable Endoscopic Tissues Reconstruction with Gaussian Splatting," in *MICCAI 2024 Workshops*. Springer, 2024, pp. 135–145.
- [41] Y. Liu, C. Li, H. Liu, C. Yang, and Y. Yuan, "Foundation Model-Guided Gaussian Splatting for 4D Reconstruction of Deformable Tissues," *IEEE Transactions on Medical Imaging*, vol. 44, no. 6, pp. 2672–2682, 2025.